

---

**QUESTION 187**

You have an on-premises network to which you deploy a virtual appliance.

You plan to deploy several Azure virtual machines and connect the on-premises network to Azure by using a Site-to-Site connection.

All network traffic that will be directed from the Azure virtual machines to a specific subnet must flow through the virtual appliance.

You need to recommend solutions to manage network traffic.

Which two options should you recommend? Each correct answer presents a complete solution.

- A. Configure Azure Traffic Manager.
- B. Implement an Azure virtual network.
- C. Implement Azure ExpressRoute.
- D. Configure a routing table.

Answer: CD

Explanation:

Connectivity can be from an any-to-any (IP VPN) network, a point-to-point Ethernet network, or a virtual cross-connection through a connectivity provider at a co-location facility. ExpressRoute connections do not go over the public Internet. This allows ExpressRoute connections to offer more reliability, faster speeds, lower latencies, and higher security than typical connections over the Internet.

Reference:

<https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-forced-tunneling-rm>

<https://docs.microsoft.com/en-us/azure/expressroute/expressroute-introduction>

---

**QUESTION 188**

You are designing a message application that will run on an on-premises Ubuntu virtual machine.

The application will use Azure Storage queues.

You need to recommend a processing solution for the application to interact with the storage queues. The solution must meet the following requirements:

Create and delete queues daily.

Be scheduled by using a CRON job.

Upload messages every five minutes.

What should developers use to interact with the queues?

- A. Azure CLI
- B. AzCopy
- C. Azure Data Factory
- D. .NET Core

Answer: D

Explanation:

Reference:

<https://docs.microsoft.com/en-us/azure/storage/queues/storage-tutorial-queues>

---

### **QUESTION** 189

#### **HOTSPOT**

You plan to deploy the backup policy shown in the following exhibit.

## Policy1

Associated Items Delete Save Discard

### Backup frequency

Daily 6:00 PM (UTC) Coordinated Universal Time

### Retention range

☒ Retention of daily backup point.

\* At For  
6:00 PM 90 Day(s)

☒ Retention of weekly backup point.

\* On \* At For  
Sunday 6:00 PM 26 Week(s)

☒ Retention of monthly backup point.

Week Based Day Based

\* On \* Day \* At For  
First Sunday 6:00 PM 36 Month(s)

☐ Retention of yearly backup point.

Not Configured

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Virtual machines that are backed up using the policy can be recovered for up to a maximum of **[answer choice]**.

	▼
90 days	
26 weeks	
36 months	
45 months	

The minimum recovery point objective (RPO) for virtual machines that are backed up by using the policy is **[answer choice]**.

	▼
1 hour	
1 day	
1 week	
1 month	
1 year	

Answer:

Virtual machines that are backed up using the policy can be recovered for up to a maximum of **[answer choice]**.

	▼
90 days	
26 weeks	
36 months	
45 months	

The minimum recovery point objective (RPO) for virtual machines that are backed up by using the policy is **[answer choice]**.

	▼
1 hour	
1 day	
1 week	
1 month	
1 year	

Explanation:

---

### QUESTION 190

#### DRAG DROP

Your company identifies the following business continuity and disaster recovery objectives for virtual machines that host sales, finance, and reporting application in the company's on-premises data center.

The finance application requires that data be retained for seven years. In the event of a disaster, the application must be able to run from Azure. The recovery in objective (RTO) is 10 minutes,

The reporting application must be able to recover point in-time data at a daily granularity. The RTO is eight hours.

The sales application must be able to fail over to second on-premises data center.

You need to recommend which Azure services meet the business continuity and disaster recovery objectives. The solution must minimize costs.

What should you recommend for each application? To answer, drag the appropriate services to the correct application. Each service may be used once or more times. You may need to drag the spin bar between panes or scroll to view content.

## Actions

Azure Backup only

Azure Site Recovery only

Azure Site Recovery and  
Azure Backup

## Answer Area



Sales:

Service or Services

Finance:

Service or Services

Reporting:

Service or Services

Answer:

Explanation:

- 1) Sales: Azure Site Recovery only
- 2) Finance: Azure Site Recovery and Azure Backup
- 3) Reporting: Azure Backup only

---

### QUESTION 191

You need to design a highly available Azure SQL database that meets the following requirements:

- \* Failover between replicas of the database must occur without any data loss.
- \* The database must remain available in the event of a zone outage.
- \* Costs must be minimized.

Which deployment option should you use?

- A. Azure SQL Database Business Critical
- B. Azure SQL Database Managed Instance Business Critical
- C. Azure SQL Database Serverless
- D. Azure SQL Database Premium

Answer: D

Explanation:

General Purpose / Standard prevents data loss through high available storage

[https://docs.microsoft.com/en-us/azure/azure-sql/database/service-tier-generalpurpose?](https://docs.microsoft.com/en-us/azure/azure-sql/database/service-tier-generalpurpose?view=azuresql)

[view=azuresql](https://docs.microsoft.com/en-us/azure/azure-sql/database/service-tier-generalpurpose?view=azuresql). This architectural model relies on high availability and reliability of Azure

Blob storage that transparently replicates database files and guarantees no data loss if underlying infrastructure failure happens. General Purpose / Standard support Zone Redundancy For General

Purpose tier the zone-redundant configuration is Generally Available in the following regions:  
<https://docs.microsoft.com/en-us/azure/azure-sql/database/high-availabilitysla?view=azuresql&tabs=azure-powershell> Without any information regarding the usage pattern, serverless is possible. Other option is D <https://docs.microsoft.com/en-us/azure/azuresql/database/serverless-tier-overview?view=azuresql>

---

### QUESTION 192

#### HOTSPOT

Your company deploys an Azure App Service Web App.

During testing the application fails under load. The application cannot handle more than 100 concurrent user sessions. You enable the Always On feature. You also configure auto-scaling to increase counts from two to 10 based on HTTP queue length.

You need to improve the performance of the application.

Which solution should you use for each application scenario? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Scenario	Solution
Store content close to end users.	<div><div>▼</div><div>Azure Redis Cache</div><div>Azure Traffic Manager</div><div>Azure Content Delivery Network</div><div>Azure Application Gateway</div></div>
Store content close to the application.	<div><div>▼</div><div>Azure Redis Cache</div><div>Azure Traffic Manager</div><div>Azure Content Delivery Network</div><div>Azure Application Gateway</div></div>

Answer:

## Scenario

## Solution

Store content close to end users.

	▼
Azure Redis Cache	
Azure Traffic Manager	
Azure Content Delivery Network	
Azure Application Gateway	

Store content close to the application.

	▼
Azure Redis Cache	
Azure Traffic Manager	
Azure Content Delivery Network	
Azure Application Gateway	

Explanation:

Box 1: Content Delivery Network

A content delivery network (CDN) is a distributed network of servers that can efficiently deliver web content to users. CDNs store cached content on edge servers in point-of-presence (POP) locations that are close to end users, to minimize latency.

Azure Content Delivery Network (CDN) offers developers a global solution for rapidly delivering highbandwidth content to users by caching their content at strategically placed physical nodes across the world. Azure CDN can also accelerate dynamic content, which cannot be cached, by leveraging various network optimizations using CDN POPs. For example, route optimization to bypass Border Gateway Protocol (BGP).

Box 2: Azure Redis Cache

Azure Cache for Redis is based on the popular software Redis. It is typically used as a cache to improve the performance and scalability of systems that rely heavily on backend data-stores.

Performance is improved by temporarily copying frequently accessed data to fast storage located close to the application. With Azure Cache for Redis, this fast storage is located in-memory with Azure Cache for Redis instead of being loaded from disk by a database.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-overview>



---

**QUESTION 193**

You plan to store data in Azure Blob storage for many years. The stored data will be accessed rarely. You need to ensure that the data in Blob storage is always available for immediate access. The solution must minimize storage costs. Which storage tier should you use?

- A. Cool
- B. Archive
- C. Hot

Answer: A

Explanation:

Azure cool tier is equivalent to the Amazon S3 Infrequent Access (S3-IA) storage in AWS that provides a low cost high performance storage for infrequently access data.

Note: Azures cool storage tier, also known as Azure cool Blob storage, is for infrequently-accessed data that needs to be stored for a minimum of 30 days. Typical use cases include backing up data before tiering to archival systems, legal data, media files, system audit information, datasets used for big data analysis and more.

The storage cost for this Azure cold storage tier is lower than that of hot storage tier. Since it is expected that the data stored in this tier will be accessed less frequently, the data access charges are high when compared to hot tier. There are no additional changes required in your applications as these tiers can be accessed using APIs in the same manner that you access Azure storage.

Reference:

<https://cloud.netapp.com/blog/low-cost-storage-options-on-azure>

---

**QUESTION 194**

HOTSPOT

You have an on-premises file server that stores 2 TB of data files.

You plan to move the data files to Azure Blob storage in the Central Europe region.

You need to recommend a storage account type to store the data files and a replication solution for the storage account. The solution must meet the following requirements:

Be available if a single Azure datacenter fails.

Support storage tiers.

Minimize cost.

What should you recommend? To answer, select the appropriate options in the answer area.  
NOTE: Each correct selection is worth one point.

Account type:

	▼
Blob storage	
Storage (general purpose v1)	
StorageV2 (general purpose v2)	

Replication solution:

	▼
Geo-redundant storage (GRS)	
Zone-redundant storage (ZRS)	
Locally-redundant storage (LRS)	
Read-access geo-redundant storage (RA-GRS)	

Answer: