

You plan to use an Azure Storage account to store data assets.

You need to recommend a solution that meets the following requirements:

- Supports immutable storage
- Disables anonymous access to the storage account
- Supports access control list (ACL)-based Azure AD permissions

What should you include in the recommendation?

- A. Azure Files
- B. Azure Data Lake Storage
- C. Azure NetApp Files
- D. Azure Blob Storage

Correct Answer: C

Community vote distribution

B (59%)

D (41%)

HOTSPOT

-

You are designing a storage solution that will ingest, store, and analyze petabytes (PBs) of structured, semi-structured, and unstructured text data. The analyzed data will be offloaded to Azure Data Lake Storage Gen2 for long-term retention.

You need to recommend a storage and analytics solution that meets the following requirements:

- Stores the processed data
- Provides interactive analytics
- Supports manual scaling, built-in autoscaling, and custom autoscaling

What should you include in the recommendation? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

For storage and interactive analytics:

▼

Azure Data Explorer
Azure Data Lake Analytics
Log Analytics

Query language:

▼

KQL
Transact-SQL
U-SQL

Answer Area

For storage and interactive analytics:

Correct Answer:

▼

Azure Data Explorer
Azure Data Lake Analytics
Log Analytics

▼

KQL
Transact-SQL
U-SQL

HOTSPOT

-

You plan to use Azure SQL as a database platform.

You need to recommend an Azure SQL product and service tier that meets the following requirements:

- Automatically scales compute resources based on the workload demand
- Provides per second billing

What should you recommend? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Azure SQL product: ▼

- A single Azure SQL database
- An Azure SQL Database elastic pool
- Azure SQL Managed Instance

Service tier: ▼

- Basic
- Business Critical
- General Purpose
- Hyperscale
- Standard

Answer Area

Correct Answer:

Azure SQL product: ▼

- A single Azure SQL database
- An Azure SQL Database elastic pool
- Azure SQL Managed Instance

Service tier: ▼

- Basic
- Business Critical
- General Purpose
- Hyperscale
- Standard

HOTSPOT

-

You have an Azure subscription.

You need to deploy a solution that will provide point-in-time restore for blobs in storage accounts that have blob versioning and blob soft delete enabled.

Which type of blob should you create, and what should you enable for the accounts? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Blob type:
Append
Block
Page

Enable:
A stored access policy
Immutable blob storage
Object replication
The change feed

Answer Area

Correct Answer:

Blob type:
Append
Block
Page

Enable:
A stored access policy
Immutable blob storage
Object replication
The change feed

HOTSPOT

-

Your company, named Contoso, Ltd., has an Azure subscription that contains the following resources:

- An Azure Synapse Analytics workspace named contosoworkspace1
- An Azure Data Lake Storage account named contosolake1
- An Azure SQL database named contososql1

The product data of Contoso is copied from contososql1 to contosolake1.

Contoso has a partner company named Fabrikam Inc. Fabrikam has an Azure subscription that contains the following resources:

- A virtual machine named FabrikamVM1 that runs Microsoft SQL Server 2019
- An Azure Storage account named fabrikamsa1

Contoso plans to upload the research data on FabrikamVM1 to contosolake1. During the upload, the research data must be transformed to the data formats used by Contoso.

The data in contosolake1 will be analyzed by using contosoworkspace1.

You need to recommend a solution that meets the following requirements:

- Upload and transform the FabrikamVM1 research data.
- Provide Fabrikam with restricted access to snapshots of the data in contosoworkspace1.

What should you recommend for each requirement? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Upload and transform the data:

Azure Data Box Gateway

Azure Data Share

Azure Synapse pipelines

Provide restricted access:

Azure Data Box Gateway

Azure Data Share

Azure Synapse pipelines

Answer Area

Upload and transform the data:

Azure Data Box Gateway

Azure Data Share

Azure Synapse pipelines

Correct Answer:

Provide restricted access:

Azure Data Box Gateway

Azure Data Share

Azure Synapse pipelines

HOTSPOT

-

You are designing a data pipeline that will integrate large amounts of data from multiple on-premises Microsoft SQL Server databases into an analytics platform in Azure. The pipeline will include the following actions:

- Database updates will be exported periodically into a staging area in Azure Blob storage.
- Data from the blob storage will be cleansed and transformed by using a highly parallelized load process.
- The transformed data will be loaded to a data warehouse.
- Each batch of updates will be used to refresh an online analytical processing (OLAP) model in a managed serving layer.
- The managed serving layer will be used by thousands of end users.

You need to implement the data warehouse and serving layers.

What should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

To implement the data warehouse:

An Apache Spark pool in Azure Synapse Analytics
An Azure Synapse Analytics dedicated SQL pool
Azure Data Lake Analytics

To implement the serving layer:

Azure Analysis Services
An Apache Spark pool Azure Synapse Analytics
An Azure Synapse Analytics dedicated SQL pool

Answer Area

To implement the data warehouse:

Correct Answer:

To implement the serving layer:

An Apache Spark pool in Azure Synapse Analytics
An Azure Synapse Analytics dedicated SQL pool
Azure Data Lake Analytics

Azure Analysis Services
An Apache Spark pool Azure Synapse Analytics
An Azure Synapse Analytics dedicated SQL pool

HOTSPOT

-

You have an Azure subscription.

You need to deploy a relational database. The solution must meet the following requirements:

- Support multiple read-only replicas.
- Automatically load balance read-only requests across all the read-only replicas.
- Minimize administrative effort

What should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Service

A single Azure SQL database
An Azure SQL Database elastic pool
Azure SQL Managed Instances

Service tier

Business Critical
Hyperscale
Premium

Answer Area

Service

A single Azure SQL database

An Azure SQL Database elastic pool

Azure SQL Managed Instances

Service tier

Business Critical

Hyperscale

Premium

Correct Answer:

You have an app named App1 that uses an Azure Blob Storage container named app1data.

App1 uploads a cumulative transaction log file named File1.txt to a block blob in app1data once every hour. File1.txt only stores transaction data from the current day.

You need to ensure that you can restore the last uploaded version of File1.txt from any day for up to 30 days after the file was overwritten. The solution must minimize storage space.

What should you include in the solution?

- A. container soft delete
- B. blob snapshots
- C. blob soft delete
- D. blob versioning

Correct Answer: D

Community vote distribution

D (67%)

B (33%)

You have 12 on-premises data sources that contain customer information and consist of Microsoft SQL Server, MySQL, and Oracle databases.

You have an Azure subscription.

You plan to create an Azure Data Lake Storage account that will consolidate the customer information for analysis and reporting.

You need to recommend a solution to automatically copy new information from the data sources to the Data Lake Storage account by using extract, transform and load (ETL). The solution must minimize administrative effort.

What should you include in the recommendation?

- A. Azure Data Factory
- B. Azure Data Explorer
- C. Azure Data Share
- D. Azure Data Studio

Correct Answer: A

Community vote distribution

A (100%)

You have SQL Server on an Azure virtual machine. The databases are written to nightly as part of a batch process.

You need to recommend a disaster recovery solution for the data. The solution must meet the following requirements:

- ☞ Provide the ability to recover in the event of a regional outage.
- ☞ Support a recovery time objective (RTO) of 15 minutes.
- ☞ Support a recovery point objective (RPO) of 24 hours.
- ☞ Support automated recovery.
- ☞ Minimize costs.

What should you include in the recommendation?

- A. Azure virtual machine availability sets
- B. Azure Disk Backup
- C. an Always On availability group
- D. Azure Site Recovery

Correct Answer: D

Replication with Azure Site Recover:

- ☞ RTO is typically less than 15 minutes.
- ☞ RPO: One hour for application consistency and five minutes for crash consistency.

Incorrect Answers:

B: Too slow.

C: Always On availability group RPO: Because replication to the secondary replica is asynchronous, there's some data loss.

Reference:

<https://docs.microsoft.com/en-us/azure/site-recovery/site-recovery-sql>

Community vote distribution

D (82%)

C (18%)

HOTSPOT -

You plan to deploy the backup policy shown in the following exhibit.

Policy 1

 Associated items  Delete  Save  Discard

Backup schedule

*Frequency *Time *Timezone

Instant Restore

Retain instant recovery snapshot(s) for


 Day(s) 

Retention range

☒ Retention of daily backup point.

*At For  Day(s)

☒ Retention of weekly backup point.

*On *At For  Week(s)

☒ Retention of monthly backup point.

Week Based

Day Based

*On *Day *At For Month(s)

☐ Retention of yearly backup point.

Not Configured

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.